

Chapter 2: Multi-armed Bandits

Ryan Lee

One-armed Bandit

- Slot machine
- Each spin (action) is independent




Multi-armed Bandit problem

- Multiple slot machines to choose from
- Simplified setting to avoid complexities of RL problems
 - No observation
 - Action does not have delayed effect



10-armed Testbed

- 10 actions, 10 reward distributions
- Reward R_t chosen from stationary probability distributions



Expected Reward

- Knowing expected reward trivializes the problem
- Estimate $q_*(a)$ with $Q_t(a)$

$$q_*(a) \doteq \mathbb{E}(R_t \mid A_t = a)$$

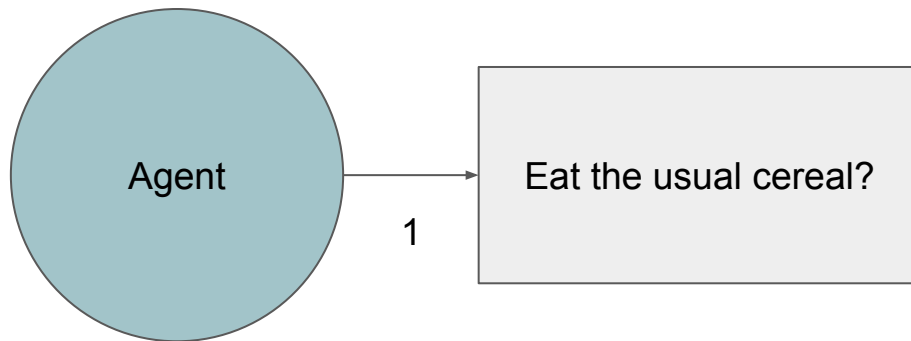
Sample-average

- Estimate $q_*(a)$ by averaging received rewards
- Default value (ex. 0) if action was never selected
- $Q_t(a)$ converges to $q_*(a)$ as denominator goes to infinity

$$Q_t(a) \doteq \frac{\text{sum of rewards when } a \text{ taken prior to } t}{\text{number of times } a \text{ taken prior to } t} = \frac{\sum_{i=1}^{t-1} R_i \cdot \mathbb{1}_{A_i=a}}{\sum_{i=1}^{t-1} \mathbb{1}_{A_i=a}}$$

Greedy method

- Always select *greedily*: $A_t \doteq \operatorname{argmax}_a Q_t(a)$
- No exploration
- Often stuck in suboptimal actions

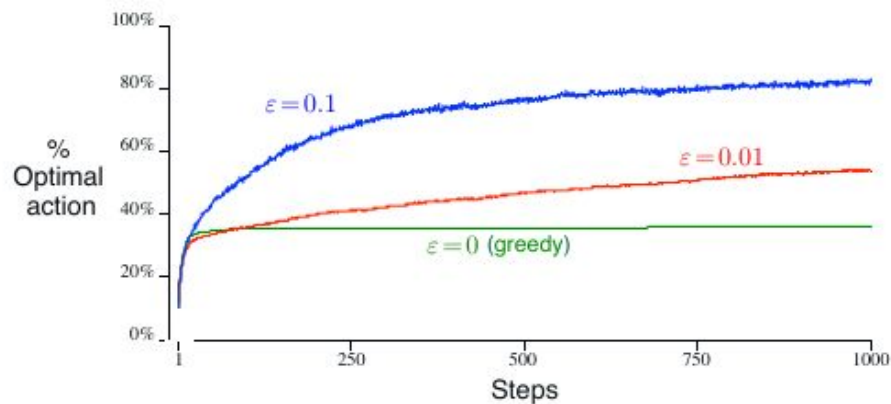
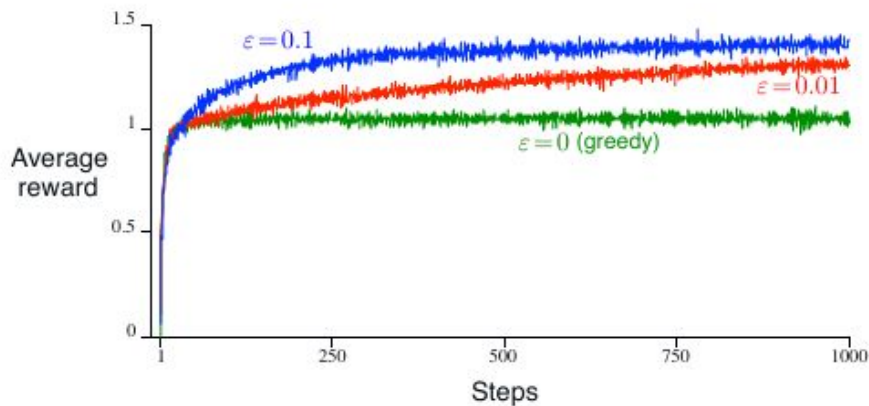


ϵ -greedy method

- Select random action with probability ϵ
- All $Q_t(a)$ converges to $q_*(a)$ as denominator goes to infinity



Greedy vs. ϵ -greedy



Incremental Implementation

- Don't store reward for each step

$$Q_{n+1} = \frac{R_1 + R_2 + \dots + R_n}{n}$$

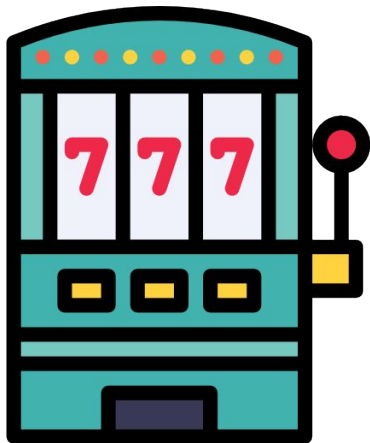
- Compute incrementally

$$Q_{n+1} = Q_n + \frac{1}{n} \left[R_n - Q_n \right]$$

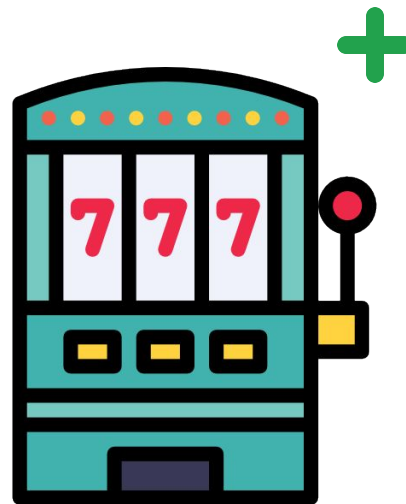
$$NewEstimate \leftarrow OldEstimate + StepSize [Target - OldEstimate]$$

Nonstationary problem

- $q_*(a)$ changes over time
- Want to give new experience more weight



$$q_*(A_1 = a) \sim N(0, 1)$$



$$q_*(A_2 = a) \sim N(3, 1)$$

Exponentially weighted average

- Constant step-size parameter α
- Give more weight to recent rewards

$$\begin{aligned} Q_{n+1} &= Q_n + \alpha [R_n - Q_n] \\ &= (1 - \alpha)^n Q_1 + \sum_{i=1}^n \alpha (1 - \alpha)^{n-i} R_i \end{aligned}$$

Sample-average

$$Q_{n+1} = Q_n + \frac{1}{n} [R_n - Q_n]$$

- Guaranteed convergence
- Converge slowly: need tuning
- Seldomly used in applications

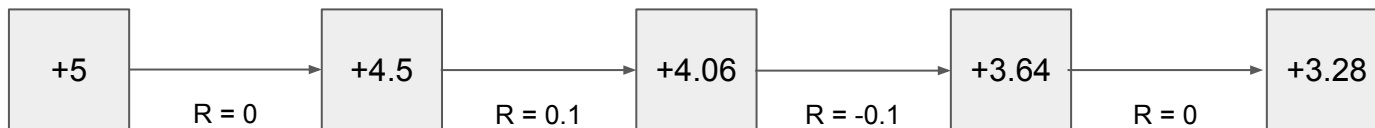
Weighted average

$$Q_{n+1} = Q_n + \alpha [R_n - Q_n]$$


- Never completely converges
- Desirable in nonstationary problems

Optimistic Initial Values

- Set initial action values optimistically (ex. +5)
- Temporarily encourage exploration
- Doesn't work in nonstationary problems



Optimistic Greedy vs. Realistic ϵ -greedy

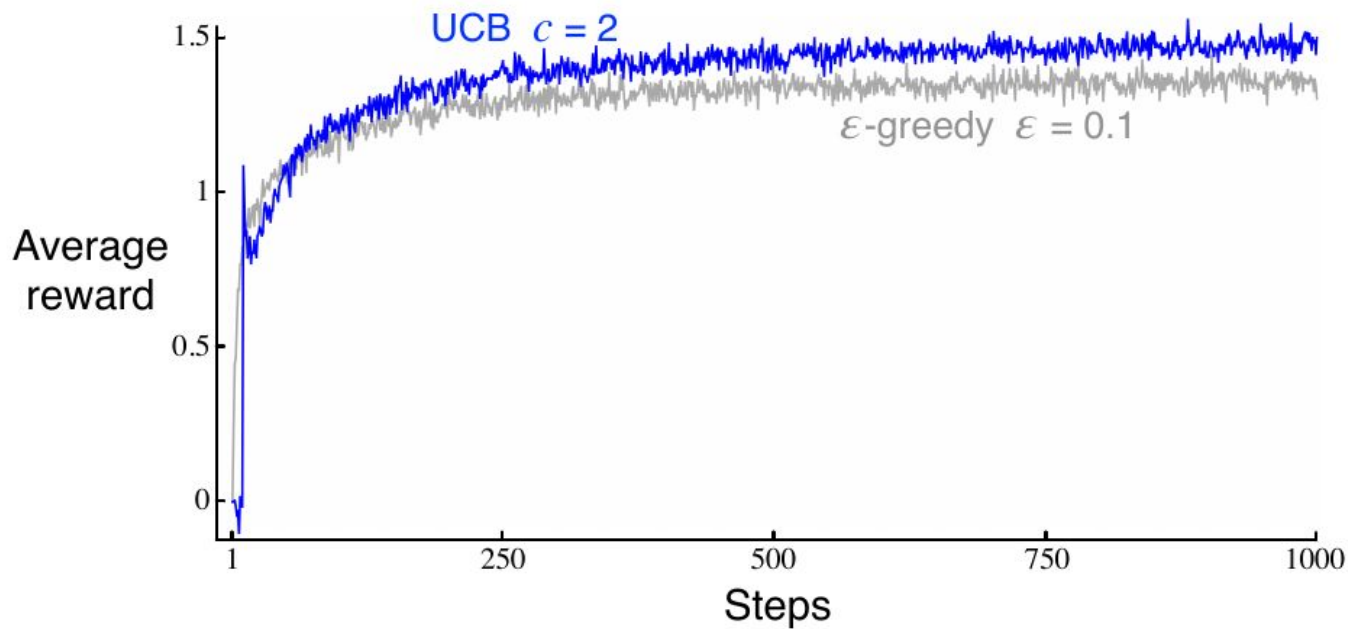


Upper Confidence Bound (UCB)

- Take into account each action's **potential** to be optimal
- Selected less \rightarrow more potential
- Difficult to extend beyond multi-armed bandits

$$A_t \doteq \operatorname{argmax} \left[Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$

UCB vs. ϵ -greedy



Gradient Bandit Algorithms

- Learn a numerical preference $H_t(a)$ for each action
- Convert to probability with softmax:

$$\pi_t(a) = \frac{e^{H_t(a)}}{\sum_{b \in \mathcal{A}} e^{H_t(b)}}$$

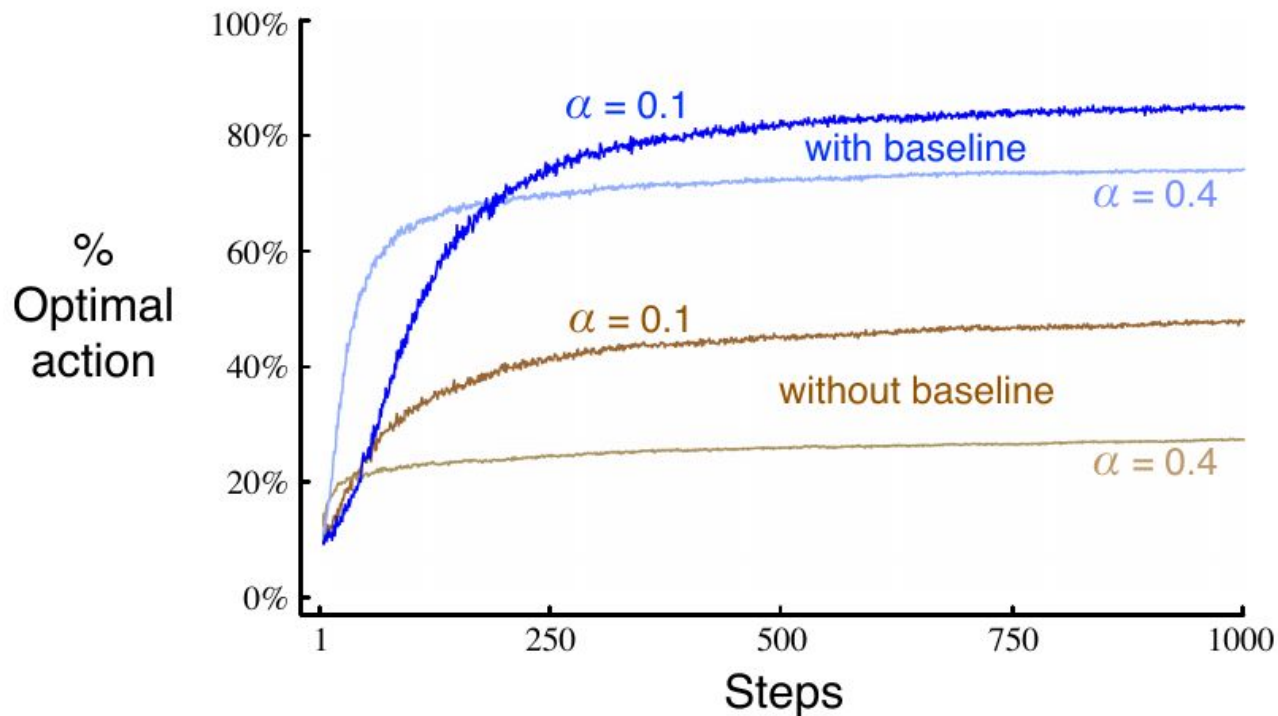
Gradient Bandit: Stochastic Gradient Descent

- Update preference $H_t(a)$ with SGD

$$\begin{aligned} H_{t+1}(A_t) &= H_t(A_t) + \alpha(R_t - \bar{R}_t)(1 - \pi_t(A_t)) \\ H_{t+1}(a) &= H_t(a) - \alpha(R_t - \bar{R}_t)\pi_t(a) \end{aligned} \quad \text{for all } a \neq A_t$$

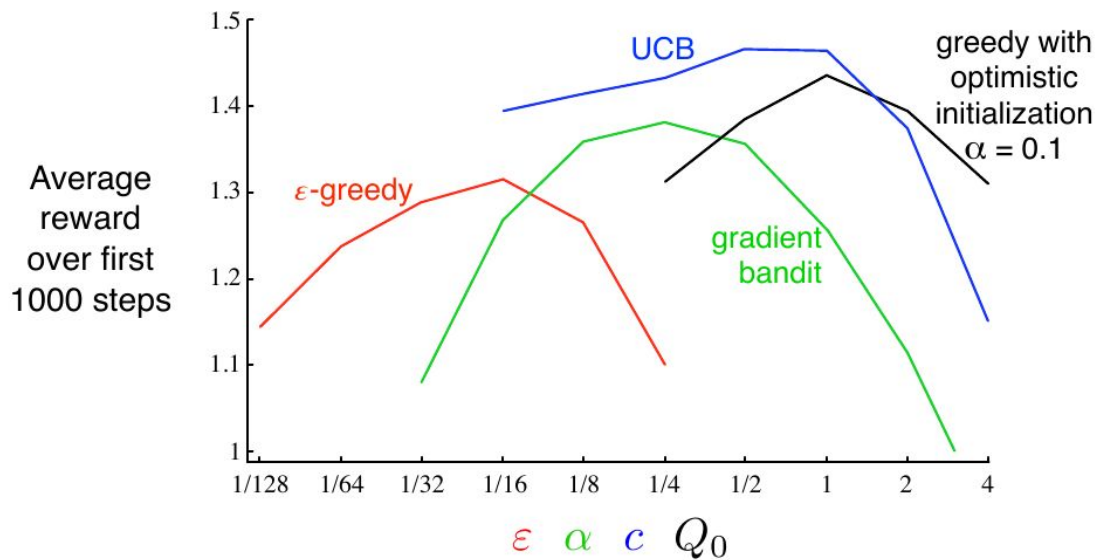
- Baseline \bar{R}_t : average of all rewards R_1, R_2, \dots, R_t
 - Increase probability if reward is above baseline
 - Decrease probability if reward is below baseline

Gradient Bandit: Results




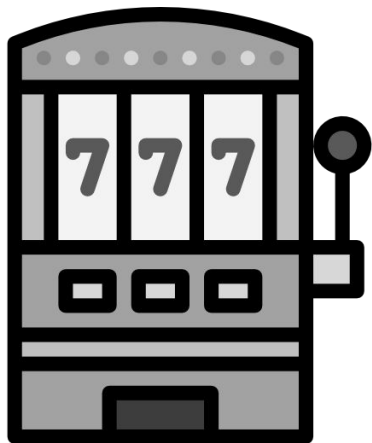
Parameter Study

- Check performance in best setting
- Check hyperparameter sensitivity



Associative Search (Contextual Bandit)

- Observe some *context* that can help decision
- Intermediate between multi-armed bandit and full RL problem
 - Need to learn a **policy** to *associate* observations and actions
 - Each action only affects immediate reward



Thank you!

Original content from

- [Reinforcement Learning: An Introduction by Sutton and Barto](#)

You can find more content in

- [github.com/seungjaeryanlee](#)
- [www.endtoend.ai](#)