

# Chapter 7: n-step Bootstrapping

Seungjae Ryan Lee

# Recap: MC vs TD

- Monte Carlo: wait until end of episode

$$V(S_t) \leftarrow V(S_t) + \alpha \left[ \overbrace{G_t - V(S_t)}^{\text{MC error}} \right],$$

Return

- 1-step TD / TD(0): wait until next time step

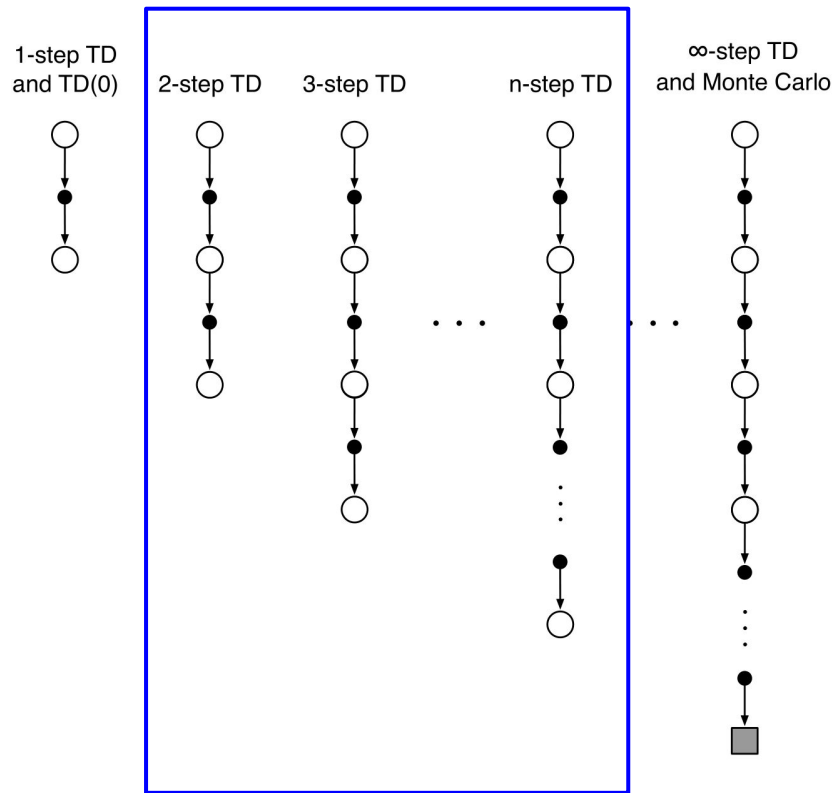
$$V(S_t) \leftarrow V(S_t) + \alpha \left[ \overbrace{R_{t+1} + \gamma V(S_{t+1}) - V(S_t)}^{\text{TD error}} \right]$$

Bootstrapping target

# n-step Bootstrapping

- Perform update based on intermediate number of rewards
- Freed from the “tyranny of the time step” of TD
  - Different time step for *action selection* (1) and *bootstrapping* interval (n)
- Called n-step TD since they still *bootstrap*

# n-step Bootstrapping



# n-step TD Prediction

- Use truncated *n-step return* as target
  - Use  $n$  rewards and bootstrap

$$G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n}),$$

- Needs future rewards not available at timestep  $t$
- $S_t$  cannot be updated until timestep  $t + n$

$$V_{t+n}(S_t) \doteq V_{t+n-1}(S_t) + \alpha [G_{t:t+n} - V_{t+n-1}(S_t)], \quad 0 \leq t < T,$$

# n-step TD Prediction: Pseudocode

*n*-step TD for estimating  $V \approx v_\pi$

Input: a policy  $\pi$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , a positive integer  $n$

Initialize  $V(s)$  arbitrarily, for all  $s \in \mathcal{S}$

All store and access operations (for  $S_t$  and  $R_t$ ) can take their index mod  $n + 1$

Loop for each episode:

    Initialize and store  $S_0 \neq \text{terminal}$

$T \leftarrow \infty$

    Loop for  $t = 0, 1, 2, \dots$ :

        If  $t < T$ , then:

            Take an action according to  $\pi(\cdot | S_t)$

            Observe and store the next reward as  $R_{t+1}$  and the next state as  $S_{t+1}$

            If  $S_{t+1}$  is terminal, then  $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$  ( $\tau$  is the time whose state's estimate is being updated)

        If  $\tau \geq 0$ :

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

            If  $\tau + n < T$ , then:  $G \leftarrow G + \gamma^n V(S_{\tau+n})$

Compute n-step return ( $G_{\tau:\tau+n}$ )

$V(S_\tau) \leftarrow V(S_\tau) + \alpha [G - V(S_\tau)]$

Update V

    Until  $\tau = T - 1$

# n-step TD Prediction: Convergence

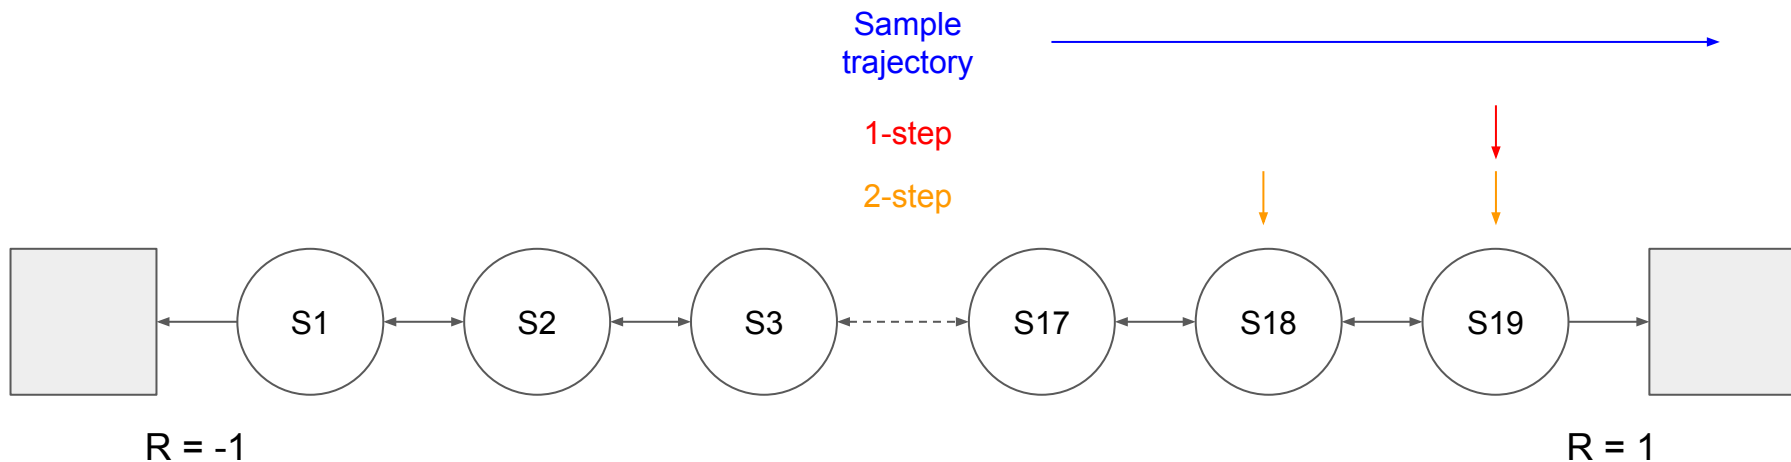
- The n-step return has the *error reduction property*
  - Expectation of n-step return is a better estimate of  $v_\pi$  than  $V_{t+n-1}$  in the worst-state sense

$$\max_s \left| \mathbb{E}_\pi[G_{t:t+n} | S_t = s] - v_\pi(s) \right| \leq \gamma^n \max_s \left| V_{t+n-1}(s) - v_\pi(s) \right|,$$

- Converges to true value under appropriate technical conditions

# Random Walk Example


- Rewards only on exit (-1 on left exit, 1 on right exit)
- n-step return: propagate reward up to n latest states





# Random Walk Example: n-step TD Prediction

- Intermediate n does best



# n-step Sarsa

- Extend n-step TD Prediction to Control (Sarsa)


- Need to use Q instead of V
- Use  $\epsilon$ -greedy policy

- Redefine n-step return with Q

$$G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q_{t+n-1}(S_{t+n}, A_{t+n}), \quad n \geq 1, 0 \leq t < T-n,$$

- Naturally extend to Sarsa


$$Q_{t+n}(S_t, A_t) \doteq Q_{t+n-1}(S_t, A_t) + \alpha [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)], \quad 0 \leq t < T,$$




# n-step Sarsa vs. Sarsa(0)

- Gridworld with nonzero reward only at the end
- n-step can learn much more from one episode


## Path taken



Action values increased by one-step Sarsa



Action values increased by 10-step Sarsa



# n-step Sarsa: Pseudocode

*n*-step Sarsa for estimating  $Q \approx q_*$  or  $q_\pi$

Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}, a \in \mathcal{A}$   
Initialize  $\pi$  to be  $\varepsilon$ -greedy with respect to  $Q$ , or to a fixed given policy  
Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$ , a positive integer  $n$   
All store and access operations (for  $S_t$ ,  $A_t$ , and  $R_t$ ) can take their index mod  $n + 1$

Loop for each episode:

- Initialize and store  $S_0 \neq \text{terminal}$
- Select and store an action  $A_0 \sim \pi(\cdot | S_0)$
- $T \leftarrow \infty$
- Loop for  $t = 0, 1, 2, \dots$ :
  - If  $t < T$ , then:
    - Take action  $A_t$
    - Observe and store the next reward as  $R_{t+1}$  and the next state as  $S_{t+1}$
    - If  $S_{t+1}$  is terminal, then:
      - $T \leftarrow t + 1$
    - else:
      - Select and store an action  $A_{t+1} \sim \pi(\cdot | S_{t+1})$
  - $\tau \leftarrow t - n + 1$  ( $\tau$  is the time whose estimate is being updated)
  - If  $\tau \geq 0$ :
    - $G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$
    - If  $\tau + n < T$ , then  $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$  ( $G_{\tau:\tau+n}$ )
    - $Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha [G - Q(S_\tau, A_\tau)]$
    - If  $\pi$  is being learned, then ensure that  $\pi(\cdot | S_\tau)$  is  $\varepsilon$ -greedy wrt  $Q$

Until  $\tau = T - 1$

# n-step Expected Sarsa


- Same update as Sarsa except the last element
  - Consider all possible actions in the last step
- Same n-step return as Sarsa except the last step

$$G_{t:t+n} \doteq R_{t+1} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n \bar{V}_{t+n-1}(S_{t+n}), \quad t+n < T,$$

$$\bar{V}_t(s) \doteq \sum_a \pi(a|s) Q_t(s, a),$$

- Same update as Sarsa

$$Q_{t+n}(S_t, A_t) \doteq Q_{t+n-1}(S_t, A_t) + \alpha [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)], \quad 0 \leq t < T,$$



# Off-policy n-step Learning

- Need *importance sampling*

$$\rho_{t:h} \doteq \prod_{k=t}^{\min(h, T-1)} \frac{\pi(A_k|S_k)}{b(A_k|S_k)}.$$

- Update target policy's values with behavior policy's returns

$$V_{t+n}(S_t) \doteq V_{t+n-1}(S_t) + \alpha \underline{\rho_{t:t+n-1}} [G_{t:t+n} - V_{t+n-1}(S_t)], \quad 0 \leq t < T,$$

- Generalizes the on-policy case
  - If  $\pi = b$ , then  $\rho = 1$

# Off-policy n-step Sarsa

- Update Q instead of V
- Importance sampling ratio starts one step later for Q values
  - $A_t$  is already chosen

$$Q_{t+n}(S_t, A_t) \doteq Q_{t+n-1}(S_t, A_t) + \alpha \underline{\rho_{t+1:t+n}} [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)],$$

# Off-policy n-step Sarsa: Pseudocode

## Off-policy $n$ -step Sarsa for estimating $Q \approx q_s$ or $q_\pi$

Input: an arbitrary behavior policy  $b$  such that  $b(a|s) > 0$ , for all  $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize  $\pi$  to be greedy with respect to  $Q$ , or as a fixed given policy

Algorithm parameters: step size  $\alpha \in (0, 1]$ , a positive integer  $n$

All store and access operations (for  $S_t$ ,  $A_t$ , and  $R_t$ ) can take their index mod  $n + 1$

Loop for each episode:

Initialize and store  $S_0 \neq$  terminal

Select and store an action  $A_0 \sim b(\cdot|S_0)$

$T \leftarrow \infty$

Loop for  $t = 0, 1, 2, \dots$ :

  If  $t < T$ , then:

    Take action  $A_t$

    Observe and store the next reward as  $R_{t+1}$  and the next state as  $S_{t+1}$

    If  $S_{t+1}$  is terminal, then:

$T \leftarrow t + 1$

    else:

      Select and store an action  $A_{t+1} \sim b(\cdot|S_{t+1})$

$\tau \leftarrow t - n + 1$  ( $\tau$  is the time whose estimate is being updated)

  If  $\tau \geq 0$ :

$$\rho \leftarrow \prod_{i=\tau+1}^{\min(\tau+n-1, T-1)} \frac{\pi(A_i|S_i)}{b(A_i|S_i)} \quad (\rho_{\tau+1:t+n-1})$$

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i \quad (G_{\tau:\tau+n})$$

  If  $\tau + n < T$ , then:  $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$

$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha \rho [G - Q(S_\tau, A_\tau)]$

  If  $\pi$  is being learned, then ensure that  $\pi(\cdot|S_\tau)$  is greedy wrt  $Q$

Until  $\tau = T - 1$



# Off-policy n-step Expected Sarsa

- Importance sampling ratio ends one step earlier for Expected Sarsa

$$Q_{t+n}(S_t, A_t) \doteq Q_{t+n-1}(S_t, A_t) + \alpha \underbrace{\rho_{t+1:t+n-1}} \left[ \underbrace{G_{t:t+n}} - Q_{t+n-1}(S_t, A_t) \right]$$

- Use expected n-step return

$$G_{t:t+n} \doteq R_{t+1} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n \bar{V}_{t+n-1}(S_{t+n}), \quad t+n < T,$$

$$\bar{V}_t(s) \doteq \sum_a \pi(a|s) Q_t(s, a),$$

# Per-decision Off-policy Methods: Intuition\*

- More efficient off-policy n-step method
- Write returns recursively:

$$G_{t:h} = R_{t+1} + \gamma G_{t+1:h},$$

$$G_{h:h} \doteq V_{h-1}(S_h).$$

- Naive importance sampling
  - If  $\rho_t = 0$  ,  $G_{t:h} = 0$
  - Estimate shrinks, higher variance

$$G_{t:h} \doteq \rho_t (R_{t+1} + \gamma G_{t+1:h})$$

$$G_{t:h} \doteq \rho_t (R_{t+1} + \gamma G_{t+1:h}) + (1 - \rho_t) V_{h-1}(S_t),$$

# Per-decision Off-policy Methods\*

- Better: If  $\rho_t = 0$  , leave the estimate unchanged

$$G_{t:h} \doteq \rho_t (R_{t+1} + \gamma G_{t+1:h}) + \underbrace{(1 - \rho_t)V_{h-1}(S_t)}_{\text{Control Variate}},$$

- Expected update is unchanged since  $\mathbb{E}[\rho_t] = 1$

$$\mathbb{E}\left[\frac{\pi(A_k|S_k)}{b(A_k|S_k)}\right] \doteq \sum_a b(a|S_k) \frac{\pi(a|S_k)}{b(a|S_k)} = \sum_a \pi(a|S_k) = 1.$$

- Used with TD update *without* importance sampling

# Per-decision Off-policy Methods: $Q^*$

- Use Expected Sarsa's n-step return

$$G_{t:t+n} \doteq R_{t+1} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n \bar{V}_{t+n-1}(S_{t+n}), \quad t+n < T,$$

$$\bar{V}_t(s) \doteq \sum_a \pi(a|s) Q_t(s, a),$$



- Off-policy form with control variate:

$$\begin{aligned} G_{t:h} &\doteq R_{t+1} + \gamma \left( \rho_{t+1} G_{t+1:h} + \bar{V}_{h-1}(S_{t+1}) - \rho_{t+1} Q_{h-1}(S_{t+1}, A_{t+1}) \right), \\ &= R_{t+1} + \gamma \rho_{t+1} \left( G_{t+1:h} - Q_{h-1}(S_{t+1}, A_{t+1}) \right) + \gamma \bar{V}_{h-1}(S_{t+1}), \quad t < h \leq T. \end{aligned}$$

- Analogous to Expected Sarsa after combining with TD update algorithm

# n-step Tree Backup Algorithm

- Off-policy *without* importance sampling
- Update from entire tree of estimated action values
  - Leaf action nodes (not selected) contribute to the target
  - Selected action nodes does not contribute but weighs all next-level action values




# n-step Tree Backup Algorithm: n-step Return

- 1-step return

$$G_{t:t+1} \doteq R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) Q_t(S_{t+1}, a),$$

- 2-step return

$$\begin{aligned} G_{t:t+2} &\doteq R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1}) Q_{t+1}(S_{t+1}, a) \\ &\quad + \gamma \pi(A_{t+1}|S_{t+1}) \left( R_{t+2} + \gamma \sum_a \pi(a|S_{t+2}) Q_{t+1}(S_{t+2}, a) \right) \\ &= R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1}) Q_{t+1}(S_{t+1}, a) + \gamma \pi(A_{t+1}|S_{t+1}) G_{t+1:t+2}, \end{aligned}$$




# n-step Tree Backup Algorithm: n-step Return

- 2-step return

$$\begin{aligned}
 G_{t:t+2} &\doteq R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1}) Q_{t+1}(S_{t+1}, a) \\
 &\quad + \gamma \pi(A_{t+1}|S_{t+1}) \left( R_{t+2} + \gamma \sum_a \pi(a|S_{t+2}) Q_{t+1}(S_{t+2}, a) \right) \\
 &= R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1}) Q_{t+1}(S_{t+1}, a) + \gamma \pi(A_{t+1}|S_{t+1}) G_{t+1:t+2},
 \end{aligned}$$

- n-step return

$$G_{t:t+n} \doteq R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1}) Q_{t+n-1}(S_{t+1}, a) + \gamma \pi(A_{t+1}|S_{t+1}) G_{t+1:t+n},$$



# n-step Tree Backup Algorithm: Pseudocode

```
Loop for each episode:
  Initialize and store  $S_0 \neq \text{terminal}$ 
  Choose an action  $A_0$  arbitrarily as a function of  $S_0$ ; Store  $A_0$ 
   $T \leftarrow \infty$ 
  Loop for  $t = 0, 1, 2, \dots$  :
    If  $t < T$ :
      Take action  $A_t$ ; observe and store the next reward and state as  $R_{t+1}, S_{t+1}$ 
      If  $S_{t+1}$  is terminal:
         $T \leftarrow t + 1$ 
      else:
        Choose an action  $A_{t+1}$  arbitrarily as a function of  $S_{t+1}$ ; Store  $A_{t+1}$ 
     $\tau \leftarrow t + 1 - n$  ( $\tau$  is the time whose estimate is being updated)
    If  $\tau \geq 0$ :
      If  $t + 1 \geq T$ :
         $G \leftarrow R_T$ 
      else
         $G \leftarrow R_{t+1} + \gamma \sum_a \pi(a|S_{t+1})Q(S_{t+1}, a)$ 
      Loop for  $k = \min(t, T - 1)$  down through  $\tau + 1$ :
         $G \leftarrow R_k + \gamma \sum_{a \neq A_k} \pi(a|S_k)Q(S_k, a) + \gamma \pi(A_k|S_k)G$ 
         $Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha [G - Q(S_\tau, A_\tau)]$ 
      If  $\pi$  is being learned, then ensure that  $\pi(\cdot|S_\tau)$  is greedy wrt  $Q$ 
    Until  $\tau = T - 1$ 
```




# A Unifying Algorithm: n-step $Q(\sigma)^*$

- Unify Sarsa, Tree Backup and Expected Sarsa
  - Decide on each step to use sample action (Sarsa) or expectation of all actions (Tree Backup)


4-step  
Sarsa




4-step  
Tree backup



4-step  
Expected Sarsa



4-step  
 $Q(\sigma)$



# A Unifying Algorithm: n-step $Q(\sigma)$ : Equations\*

- $\sigma_t \in [0, 1]$ : degree of sampling on timestep  $t$

$$G_{t:h} \doteq R_{t+1} + \gamma \left( \underbrace{\sigma_{t+1} \rho_{t+1} + (1 - \sigma_{t+1}) \pi(A_{t+1} | S_{t+1})}_{\text{importance sampling}} \right) \left( G_{t+1:h} - Q_{h-1}(S_{t+1}, A_{t+1}) \right) + \gamma \bar{V}_{h-1}(S_{t+1}), \quad (7.17)$$

- Slide linearly between two weights:
  - Sarsa: Importance sampling ratio  $\rho_{t+1}$
  - Tree Backup: Policy probability  $\pi(A_{t+1} | S_{t+1})$

# A Unifying Algorithm: n-step $Q(\sigma)$ : Pseudocode\*


```
Loop for each episode:
  Initialize and store  $S_0 \neq \text{terminal}$ 
  Choose and store an action  $A_0 \sim b(\cdot|S_0)$ 
   $T \leftarrow \infty$ 
  Loop for  $t = 0, 1, 2, \dots$ :
    If  $t < T$ :
      Take action  $A_t$ ; observe and store the next reward and state as  $R_{t+1}, S_{t+1}$ 
      If  $S_{t+1}$  is terminal:
         $T \leftarrow t + 1$ 
      else:
        Choose and store an action  $A_{t+1} \sim b(\cdot|S_{t+1})$ 
        Select and store  $\sigma_{t+1}$ 
        Store  $\frac{\pi(A_{t+1}|S_{t+1})}{b(A_{t+1}|S_{t+1})}$  as  $\rho_{t+1}$ 
     $\tau \leftarrow t - n + 1$  ( $\tau$  is the time whose estimate is being updated)
    If  $\tau \geq 0$ :
       $G \leftarrow 0$ 
      Loop for  $k = \min(t + 1, T)$  down through  $\tau + 1$ :
        if  $k = T$ :
           $G \leftarrow R_T$ 
        else:
           $\bar{V} \leftarrow \sum_a \pi(a|S_k)Q(S_k, a)$ 
           $G \leftarrow R_k + \gamma(\sigma_k \rho_k + (1 - \sigma_k)\pi(A_k|S_k))(G - Q(S_k, A_k)) + \gamma \bar{V}$ 
           $Q(S_k, A_k) \leftarrow Q(S_k, A_k) + \alpha [G - Q(S_k, A_k)]$ 
      If  $\pi$  is being learned, then ensure that  $\pi(\cdot|S_\tau)$  is greedy wrt  $Q$ 
    Until  $\tau = T - 1$ 
```

# Summary


- n-step: Look ahead to the next  $n$  rewards, states, and actions
  - + Perform better than either MC or TD
  - + Escapes the *tyranny of the single time step*
  - Delay of  $n$  steps before learning
  - More memory and computation per timestep
- Extended to Eligibility Traces (Ch. 12)
  - + Minimize additional memory and computation
  - More complex
- Two approaches to off-policy n-step learning
  - Importance sampling: high variance
  - Tree backup: limited to few-step bootstrapping if policies are very different (even if  $n$  is large)

# Summary


4-step  
Sarsa




4-step  
Tree backup



4-step  
Expected Sarsa



4-step  
 $Q(\sigma)$



# Thank you!

Original content from

- [Reinforcement Learning: An Introduction by Sutton and Barto](#)

You can find more content in

- [github.com/seungjaeryanlee](#)
- [www.endtoend.ai](#)